

AI SYNERGY HACKATHON 2026

Consolidated Rules & Synergy Guidebook

Organizer-review draft 1.0 · 28 September 2026

Find the problem. Defend your reasoning. Build the solution.

We assess how you think with AI, what you choose to build, and how you respond when evidence or constraints change.

1. Purpose and format

Organized by Sangillence, in partnership with Gwalior AI Summit 2026 at ABV-IIITM Gwalior, the AI Synergy Hackathon is a two-stage challenge. Teams investigate a case, identify a specific problem, propose a scoped solution and demonstrate a working implementation.

Stage 1 — Think: Online case investigation, problem framing and a basic Software Requirements Specification (SRS), supported by genuine AI interaction logs.

Stage 2 — Build: On-campus screening, live development, a new constraint, and a final demonstration for qualifying teams.

2. Seven themes

- AI for All
- AI in Agriculture
- AI in Healthcare
- AI in Tourism
- AI for Energy & Environment
- AI in Sports
- AI in Governance

Choose one case within one theme. A theme is a domain; a case is the situation to investigate; your problem statement is the specific failure, unmet need or uncertainty your team chooses to address. Different well-supported problem statements may emerge from the same case.

3. Key information

Team size: 2–4 members · **Fee:** ₹100 per participant, non-refundable.

Registration deadline: 12 October 2026.

Stage 1 submission deadline: 12 October 2026.

Stage 2 venue: ABV-IIITM Gwalior; date and duration to be announced.

Prize pool: ₹1,00,000 target, subject to sponsorship confirmation.

Draft status: the closing times in IST, submission portal, Stage 2 schedule and screening format are not confirmed in the supplied guidebook. New operational provisions and scoring weights in this consolidated draft are identified on the organizer checklist on page 8. This draft is not a registration announcement.

Investigate before proposing.

4. Eligibility and team rules

- Open to currently enrolled undergraduate and postgraduate students from any discipline. Prior coding or AI/ML experience is not required.
- Each participant may belong to exactly one team. Cross-college teams are allowed. Duplicate participation is grounds for disqualification.
- Teams contain 2–4 members and are locked when Stage 1 submission opens. No merging, splitting or member substitution after the lock.
- Team names must not contain trademarked terms, offensive language or misleading references, as stated in the supplied rules.
- Nominate one team contact for registration, submission and official communication. Every member remains responsible for submission integrity.

5. Registration and case access

After registration, teams receive the seven case spaces, this guidebook, a common framing prompt and submission instructions. The final publication must specify the exact registration cutoff, Stage 1 opening time and submission cutoff in IST. Teams should not assume registration on the deadline date grants extra submission time.

6. Stage 1 — Required submission

Submit one package per team. The main brief is limited to six pages, excluding references, log files and the log index. Use PDF for the brief. Additional formats and upload limits will be listed with the portal instructions.

- **Problem chosen:** Identify the affected user, the precise failure or unmet need, supporting evidence and why this problem deserves attention. Compare at least two alternative problem framings before selecting your focus.
- **Evidence and assumptions:** Distinguish documented facts, stakeholder claims, inferences, assumptions and unknowns. Cite sources and state what evidence would change your decision.
- **Proposed solution / basic SRS:** Define users, inputs, outputs, the core workflow, functional requirements, relevant constraints, AI and human responsibilities, error or missing-data handling, three acceptance tests and version-1 exclusions.
- **Evaluation:** Explain how your intervention will be compared with the current workflow or a simple alternative. Include an outcome measure and a possible harm or exclusion measure.
- **Full AI logs:** Submit genuine task-related interactions across the AI tools used, plus an index pointing to consequential decisions. Include a model/tool list, contributor attribution and a declaration of reused sources or assets.

7. Qualification

Stage 1 scores determine qualification for Stage 2. Shortlist capacity and any allocation of places across themes must be published before submissions open. Qualified teams receive campus instructions. Non-qualifying teams remain eligible for the Stage 1 recognition described in Section 17.

A polished presentation cannot substitute for evidence. There is no mandatory technology stack, single correct diagnosis or requirement to place AI in every part of the proposed solution.

Build, test and adapt.

8. Screening and ownership check

The screening exercise verifies that members understand their Stage 1 investigation and can explain their contribution. Questions may refer to submitted logs, source choices, assumptions and SRS decisions. It is an understanding and attribution check, not an undisclosed domain-trivia test. Its exact format, duration and any effect on progression will be announced to qualifiers in advance.

9. Live build

Teams develop or extend a working slice of their Stage 1 proposal during the announced build period. The initial artifact and reused components must be declared at check-in. Judges assess the work done during the permitted period and the submitted prototype's behavior.

Code and no-code tools are permitted. A qualifying working slice must accept inputs and demonstrate the relevant workflow or output. Static slides, mockup screens or a video alone do not satisfy the execution requirement. Simulated integrations are allowed when clearly labeled; they do not demonstrate real deployment or measured field impact.

10. New evidence or constraint

Organizers introduce a new constraint or evidence item at a published point in the live round. Every team in the same case receives the same item at the same elapsed time. Before release, teams preserve their initial approach and artifact checkpoint.

Submit a concise revision note: what you initially believed; what changed; what you retained; the resulting requirement change; and how you tested the response. A decision to retain the original approach can receive full credit if supported by the new evidence. Change for its own sake earns no advantage.

11. Final submission and demonstration

- A runnable artifact, with a fixed submitted version or timestamped export; source code or an editable no-code export where available; and clear run instructions.
- Updated SRS or a change note, baseline comparison, test results and known limitations.
- Genuine Stage 2 AI interaction logs, contributor notes and a list of generated, reused and newly implemented components.
- A live demonstration showing the ordinary workflow, missing or conflicting information, an adverse case and the effect of the new constraint.
- A clear explanation of what AI produced, what the team decided and what was actually verified.

12. Timing and technical failures

The final event notice will state build duration, constraint-release time, artifact-freeze time and presentation/Q&A; allowances. After the freeze, teams may run the submitted version but may not alter it except under an organizer-authorized recovery procedure.

Declare unavailable external services and any mock used in their place. A confirmed shared infrastructure failure may trigger a common recovery arrangement. Private extensions or different undisclosed constraints must not advantage individual teams.

Make the process verifiable.

13. Permitted AI use and assistance

AI may support investigation, research, reasoning, design, coding, analysis, testing and communication. Teams choose the tools and remain accountable for their outputs. Paid-tool use earns no independent credit; judges evaluate the work and evidence rather than a tool's brand or cost.

Public references, open-source libraries and general-purpose templates may be reused with attribution and applicable permissions. Case-specific solutions substantially built before Stage 1 opens are ineligible. Declare the origin and extent of all pre-existing assets. External mentors may explain concepts or give feedback; they may not complete the team's assessed work or act as undeclared team members.

14. AI logs and attribution

- Preserve original task-related sessions, including failed attempts and changes of direction. Do not rewrite prompts, remove inconvenient responses, fabricate interaction histories or submit a retrospective summary as an original log.
- Use native exports or accessible share links. For tools without export support, preserve chronological transcripts or captures with enough context to verify the interaction. Disclose capture gaps and their reason.
- The log index should identify the tool/model, session, contributor or participating members, and the decisions it supports. Clearly distinguish collaborative sessions from individual work.
- Submit only material relevant to the event. Do not include unrelated private conversations, passwords, API keys or private third-party records.
- If sensitive information was accidentally included, mark its exact location and use the organizer's confidential redaction procedure. Redactions must be declared and must not alter the reasoning record. Organizers must publish that procedure before collection.

Reasoned acceptance of a verified AI answer can demonstrate strong collaboration. Judges must not require manufactured disagreement, excessive prompting or a deliberately flawed first answer.

15. Submission handling

The portal instructions govern naming, file limits and accepted log formats. The team contact must ensure links and files are accessible. If a submitted link fails, organizers may request access to the same timestamped material; this is not an opportunity to add work or revise the submission.

Late submissions are not accepted unless organizers announce a common extension. Missing required components may make a submission ineligible. Suspected fabrication or plagiarism must be reviewed against evidence, with the team given an opportunity to explain before a final integrity decision.

16. Case facts and prototype boundaries

Do not present synthetic records as field evidence, simulation results as real impact or model confidence as verified truth. Label assumptions and fabricated demo fixtures. Use public or authorized data. Where a case involves health, children or infrastructure, keep prototypes within its stated simulation and human-review boundaries; no live intervention is required.

Score the decisions and the build.

The following weights form the recommended scoring system for this consolidated draft. Publish the adopted weights before Stage 1 opens; do not change the ranking formula after teams submit. Scores for each stage are out of 100.

Stage 1 — Think

Criterion	Points
Problem diagnosis and choice: precise failure, alternatives and justified focus	25
Evidence quality: sources, uncertainty and competing explanations	20
AI collaboration: purposeful use, verification and consequential decisions	25
SRS quality: feasible scope, workflow, responsibilities and testability	20
Context and accountability: practical constraints, inclusion and recourse	10

Stage 2 — Build

Criterion	Points
Execution: runnable behavior against declared and adverse tests	30
Adaptation: justified response to the new evidence or constraint	20
AI orchestration: task allocation, corrections and verification during build	20
Evaluation and safeguards: baseline, outcomes, uncertainty and limits	20
Communication and ownership: clear explanation and team understanding	10

Finalist ranking

Final score = 30% of Stage 1 score + 70% of Stage 2 score.

Stage 1 also determines qualification. ARI feedback does not add a separate undisclosed scoring layer.

Tie-break order: Stage 2 execution score; Stage 2 adaptation score; Stage 1 evidence-quality score. If still tied, use a common short scenario with the same scoring anchors. Do not use prompt count, spending, submission speed or model brand as hidden tie-breaks.

Common scoring anchors

0: absent · **1:** asserted without support · **2:** plausible but incomplete · **3:** traceable and adequately tested · **4:** tested against alternatives and limitations. Criterion points = criterion weight × anchor ÷ 4.

Use two independent judges where staffing permits. A difference greater than one anchor on a criterion should be reconciled against cited evidence. Calibrate panels across themes; an unfamiliar but well-supported solution is not a wrong answer.

Eight dimensions. Observable evidence.

Sangillence's AI Readiness Index (ARI) guides observation and feedback across both stages. It examines how participants work with AI in this task. It does not assume that polished output proves understanding or that a short chat proves low ability.

AI Literacy

Recognizes unsupported, hallucinated, biased or overconfident output; checks important claims and understands tool limitations.

Prompt Reasoning

Breaks the task into useful questions, supplies relevant context and changes the prompting strategy when it stops helping.

Augmentation Instinct

Decides what to delegate to AI, what can be handled with simpler methods and what requires human responsibility.

Error Correction

Detects, investigates and repairs consequential errors. Does not invent mistakes merely to demonstrate correction.

Contextual Synthesis

Uses case-specific infrastructure, economics, geography, culture and user needs to change requirements or decisions.

Creative Divergence

Explores meaningfully different approaches and explains why the chosen one fits. Novelty is assessed through the work, not a claim that AI could never have proposed it.

Ethics Sensitivity

Turns concerns about consent, privacy, exclusion, fairness and accountability into concrete design decisions.

Collaborative Agency

Maintains ownership of goals, directs iterations, validates outputs and explains the final decisions.

Synergy Reports and individual attribution

Reports describe observed collaboration and provide development feedback. Team-level evidence must be labeled as team-level evidence. An individual report may include that member's attributable contributions and screening observations; it must not silently assign a shared conversation's behavior to every member.

Mark dimensions with inadequate evidence as "insufficient evidence." A report is feedback on this challenge, not a clinical assessment, fixed personality label or comprehensive measure of a participant's capability. No participant needs to produce artificial interaction volume to obtain a rating.

Clear commitments. Fair participation.

17. Prizes and recognition

Stage 1 submitters: participation certificates and individual Synergy Reports, with evidence-attribution limits as described in the ARI guidance.

Stage 2 qualifying teams: PPI opportunities, a presentation slot/campus visit at ABV-IIITM Gwalior and goodies.

Stage 2 winners: a share of the announced prize pool.

The ₹1,00,000 pool is a sponsorship-dependent target. If sponsorship falls short, the revised prize structure must be communicated before Stage 2 begins. The supplied guidebook guarantees certificates, Synergy Reports, PPI access and the qualifying-team campus visit independently of sponsorship; this draft retains those commitments. Organizers must verify delivery arrangements before release.

A PPI opportunity is an interview opportunity, not a job or internship offer. Partner-specific eligibility, scheduling and selection conditions must be communicated clearly. No travel or accommodation reimbursement is promised by this rulebook; any support must be separately announced.

18. Conduct and integrity review

Participants, mentors, judges and volunteers must treat others respectfully and refrain from harassment, intimidation and discrimination. Fabrication, plagiarism, impersonation, deliberate disruption, unauthorized access or tampering with another team's work are prohibited.

Judges and mentors must disclose conflicts of interest and avoid evaluating teams where impartiality is compromised. Organizers must identify a confidential contact for conduct complaints and a channel for reporting factual or arithmetic scoring errors.

Organizers may seek clarification, correct an administrative error, remove a participant or disqualify a submission according to the severity and evidence. Disqualification or removal for misconduct may result in forfeiture of prizes and benefits without a fee refund, as in the supplied guidebook. Record and communicate the reason to the affected team.

19. Intellectual property and public use

Teams retain ownership of their original ideas and submissions, subject to any third-party rights in reused material. Attribution and license obligations continue to apply.

The supplied guidebook grants organizers a non-exclusive, royalty-free license for event-related promotional and archival use of submission materials. This consolidated draft narrows public use to the submitted brief, build and presentation materials designated for publicity. Private AI logs and identifiable Synergy Reports require separate permission before public release. This privacy clarification is an amendment for organizer adoption before registration.

20. Access, clarifications and changes

Evaluation materials are accessible only to authorized organizers, reviewers and service providers needed for evaluation. Before collection, publish the access arrangements, retention period, deletion/contact process and any AI-assisted evaluation disclosure.

Material clarifications must be shared with all affected teams through one official channel. Publish one dated, versioned rulebook and one clarification register. Do not apply undisclosed retrospective scoring changes. Explain any unavoidable schedule or operational change and provide an equitable response.

Close these items before launch.

This page records what was preserved, what this draft adds and what still needs an operational decision. Remove it from the participant edition once the public rules and outstanding event details are finalized.

Preserved from the supplied guidebook

Sangillence / Gwalior AI Summit / ABV-IIITM context; current-college-student eligibility; 2-4 member teams; cross-college participation; ₹100 per participant; non-refundable fee; 12 October 2026 registration and Stage 1 dates; online Think and campus Build stages; required problem statement, SRS and AI logs; eight ARI dimensions; sponsor-dependent prize target; stated recognition; team ownership of original work.

Consolidated or added in this draft

- Seven themes, including Governance, aligned with the theme image and the revised case-study booklet.
- Six-page Stage 1 brief; candidate-problem comparison; evidence labels; log index, contributor attribution and baseline evaluation.
- Runnable code or no-code scope; declared mock integrations; checkpoint and revision note; common constraint-release and recovery arrangements.
- 100-point rubrics; 30/70 finalist weighting; scoring anchors, tie-breaks and judge reconciliation.
- Original-log preservation with a declared confidential redaction procedure; attribution limits on individual reports.
- Clarification of pre-existing components and mentor assistance; narrowed public-use permission for private logs and identifiable reports.

Details still to fill — no values invented

- **Deadlines:** exact registration and submission cutoff times in IST; Stage 1 opening time. Both dates are currently 12 October. Decide whether same-day late registration gives teams adequate time, and publish any revised date prospectively.
- **Stage 2:** confirmed date, build duration, screening format and its progression effect; presentation/Q&A; duration; artifact freeze and constraint-release timing.
- **Qualification:** total capacity and whether places are allocated by theme or overall score.
- **Operations:** portal, accepted files and size limits, log-capture instructions, official contact and clarification channel, outage arrangements and complaints route.
- **Benefits:** prize sponsorship and distribution; PPI partner commitments and eligibility; certificate/report delivery dates; goodies and campus arrangements.
- **Evaluation data:** common prototype fixtures, review access, permitted evaluation services, retention period, redaction and deletion process.

Release check

Resolve the fields above, adopt or revise the added provisions, and publish the final version before submissions open. Do not send private constraint cards or organizer answer guidance with the public case booklet. Pilot at least one submission and one demo against the rubric to catch ambiguous rules.

Editorial basis: the user-supplied Rules, Regulations & Synergy Guidebook and seven-theme image. This is a consolidated event-policy draft; it has not been published or sent to participants.